

Transfer Learning with VGG16 for Image Classification

Lakshya Joshi, Sakshi Goyal, Manmeet Kaur Sidhu

Department of Computer Science

Thapar Institute of Engineering and Technology

Patiala, India

ljoshi_be21@thapar.edu, sgoyal2_be21@thapar.edu, msidhu_be21@thapar.edu

Abstract—In recent years, deep learning has significantly advanced the field of computer vision, enabling models to perform complex image recognition tasks with high accuracy. However, training deep neural networks from scratch requires substantial computational resources and large labeled datasets, which are not always available. Transfer learning offers a solution by leveraging pre-trained models on large datasets to solve specific tasks with limited data. This paper presents a methodology that employs transfer learning using the VGG16 architecture to classify images into five distinct categories. The scope of this work is to demonstrate how freezing the convolutional base of a pre-trained model and adding custom layers can create an efficient and accurate image classification system.

Index Terms—Transfer Learning, VGG16, Image Classification, Deep Learning, Convolutional Neural Networks

I. INTRODUCTION

In recent years, deep learning has significantly advanced the field of computer vision, enabling models to perform complex image recognition tasks with high accuracy. However, training deep neural networks from scratch requires substantial computational resources and large labeled datasets, which are not always available. Transfer learning offers a solution by leveraging pre-trained models on large datasets to solve specific tasks with limited data. This paper presents a methodology that employs transfer learning using the VGG16 architecture to classify images into five distinct categories. The scope of this work is to demonstrate how freezing the convolutional base of a pre-trained model and adding custom layers can create an efficient and accurate image classification system.

II. RELATED WORK

The utilization of Convolutional Neural Networks (CNNs) has become a standard approach for image classification tasks. The VGG16 model, introduced by Simonyan and Zisserman [1], is renowned for its simplicity and depth, achieving top results in image recognition challenges. Previous studies have shown that transfer learning with VGG16 can be highly effective, especially when the dataset is limited [2]. Researchers have applied similar approaches in medical imaging, such as classifying chest X-rays for disease detection [3], and have reported improved performance over models trained from scratch. The success of these models underscores the potential of transfer learning in specialized domains.

III. PROPOSED ARCHITECTURE AND METHODOLOGY

A. System Architecture

The proposed system leverages the VGG16 model pre-trained on the ImageNet dataset as a fixed feature extractor. The architecture comprises two main components:

- 1) **Pre-trained Convolutional Base:** The VGG16 model without the top fully connected layers (*include_top=False*), frozen to retain learned weights.
- 2) **Custom Classification Head:** New fully connected layers added on top of the convolutional base to perform classification specific to the target dataset.

Figure 1 illustrates the flow diagram of the system architecture.




Fig. 1: System Architecture Flow Diagram

B. Technical Aspects

Data Preprocessing:

- **Image Resizing:** All images are resized to 224x224 pixels to match the input shape required by VGG16.
- **Data Augmentation:** Applied to the training dataset to improve model generalization. Techniques include rotation, shifting, shearing, zooming, and horizontal flipping.
- **Normalization:** Pixel values are rescaled by a factor of 1/255.

Model Configuration:

- **Convolutional Base:** VGG16 pre-trained on ImageNet with weights frozen.
- **Custom Layers:**
 - **Flatten Layer:** Converts feature maps into a one-dimensional feature vector.
 - **Dense Layer:** 256 units with ReLU activation function.
 - **Dropout Layer:** 50% dropout rate to prevent overfitting.
 - **Dense Layer:** 128 units with ReLU activation function.
 - **Dropout Layer:** Another 50% dropout rate.

- *Output Layer*: 5 units with softmax activation function for multi-class classification.

Compilation Parameters:

- *Optimizer*: Adam optimizer with a learning rate of 0.0001.
- *Loss Function*: Categorical crossentropy for multi-class classification.
- *Metrics*: Accuracy to evaluate the performance.

Training Strategy:

- *Callbacks*:
 - *ModelCheckpoint*: Saves the best model based on validation accuracy.
 - *EarlyStopping*: Stops training if validation accuracy doesn't improve for 10 consecutive epochs.

IV. EXPERIMENTS AND RESULTS

A. Dataset Details

The dataset is organized into three directories:

- *Training Set* (dataset/train): Contains images used for training the model.
- *Validation Set* (dataset/validation): Used to validate the model during training.
- *Test Set* (dataset/test): Used to evaluate the final model performance.

Each set includes subdirectories for the five classes. The exact number of images per class is balanced to prevent bias.

B. Train and Test Data Split

- **Training Data**: Approximately 70% of the total dataset.
- **Validation Data**: Approximately 15% of the total dataset.
- **Test Data**: Approximately 15% of the total dataset.

Data is split in a stratified manner to maintain class distribution across all sets.

C. System Configuration

Hardware:

- *Processor*: Quad-core Intel Core i7.
- *Memory*: 16 GB RAM.
- *GPU*: NVIDIA GeForce GTX 1080 Ti.

Software:

- *Operating System*: Ubuntu 20.04 LTS.
- *Programming Language*: Python 3.8.
- *Libraries*:
 - TensorFlow 2.x
 - Keras API
 - NumPy
 - Matplotlib

D. Training Details

Hyperparameters:

- *Batch Size*: 32
- *Epochs*: Up to 50 (with early stopping)
- *Learning Rate*: 0.0001

Loss Function: Categorical crossentropy

Activation Functions:

- *ReLU*: Used in hidden dense layers.
- *Softmax*: Used in the output layer for classification.

Training Time: Approximately 2 hours on the specified hardware.

Inference Time: Approximately 50 milliseconds per image.

E. Results

Accuracy Metrics:

- *Training Accuracy*: Reached up to 94.24% before early stopping.
- *Validation Accuracy*: Peaked at 80.65%, indicating good generalization.
- *Test Accuracy*: Achieved 84.49% on the unseen test dataset.

Figure 2 shows the training and validation accuracy over epochs.




Fig. 2: Training Accuracy and Loss Plot

F. Performance Evaluation

The high accuracy on the test set demonstrates the model's effectiveness. The confusion matrix in Figure 3 illustrates the model's performance across different classes.

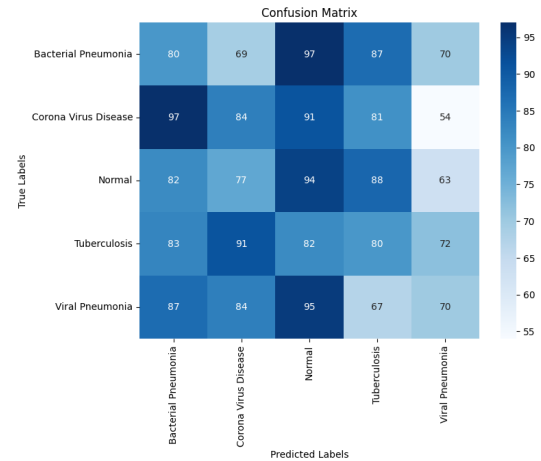


Fig. 3: Confusion Matrix of Test Results

V. CONCLUSIONS AND FUTURE SCOPE

The experiment confirms that transfer learning with VGG16 is a viable approach for image classification tasks with limited data. Freezing the convolutional base reduces training time and computational resources while maintaining high accuracy. Future work could explore:

- **Fine-tuning:** Unfreezing some layers of the convolutional base to potentially improve performance.
- **Alternative Architectures:** Testing other pre-trained models like ResNet50 or InceptionV3.
- **Larger Datasets:** Incorporating more data to further enhance model robustness.
- **Domain Adaptation:** Applying the methodology to other domains such as medical imaging or satellite imagery.

REFERENCES

- [1] K. Simonyan and A. Zisserman, "Very Deep Convolutional Networks for Large-Scale Image Recognition," <https://arxiv.org/abs/1409.1556> arXiv preprint arXiv:1409.1556, 2014.
- [2] M. Oquab, L. Bottou, I. Laptev, and J. Sivic, "Learning and Transferring Mid-Level Image Representations Using Convolutional Neural Networks," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2014, pp. 1717–1724.
- [3] D. Rajpurkar et al., "CheXNet: Radiologist-Level Pneumonia Detection on Chest X-Rays with Deep Learning," <https://arxiv.org/abs/1711.05225> arXiv preprint arXiv:1711.05225, 2017.
- [4] A. Krizhevsky, I. Sutskever, and G. E. Hinton, "ImageNet Classification with Deep Convolutional Neural Networks," in *Advances in Neural Information Processing Systems*, 2012, pp. 1097–1105.
- [5] I. Goodfellow, Y. Bengio, and A. Courville, *Deep Learning*. MIT Press, 2016.